Docker

단계1: Docker Server 실행 및 로그인

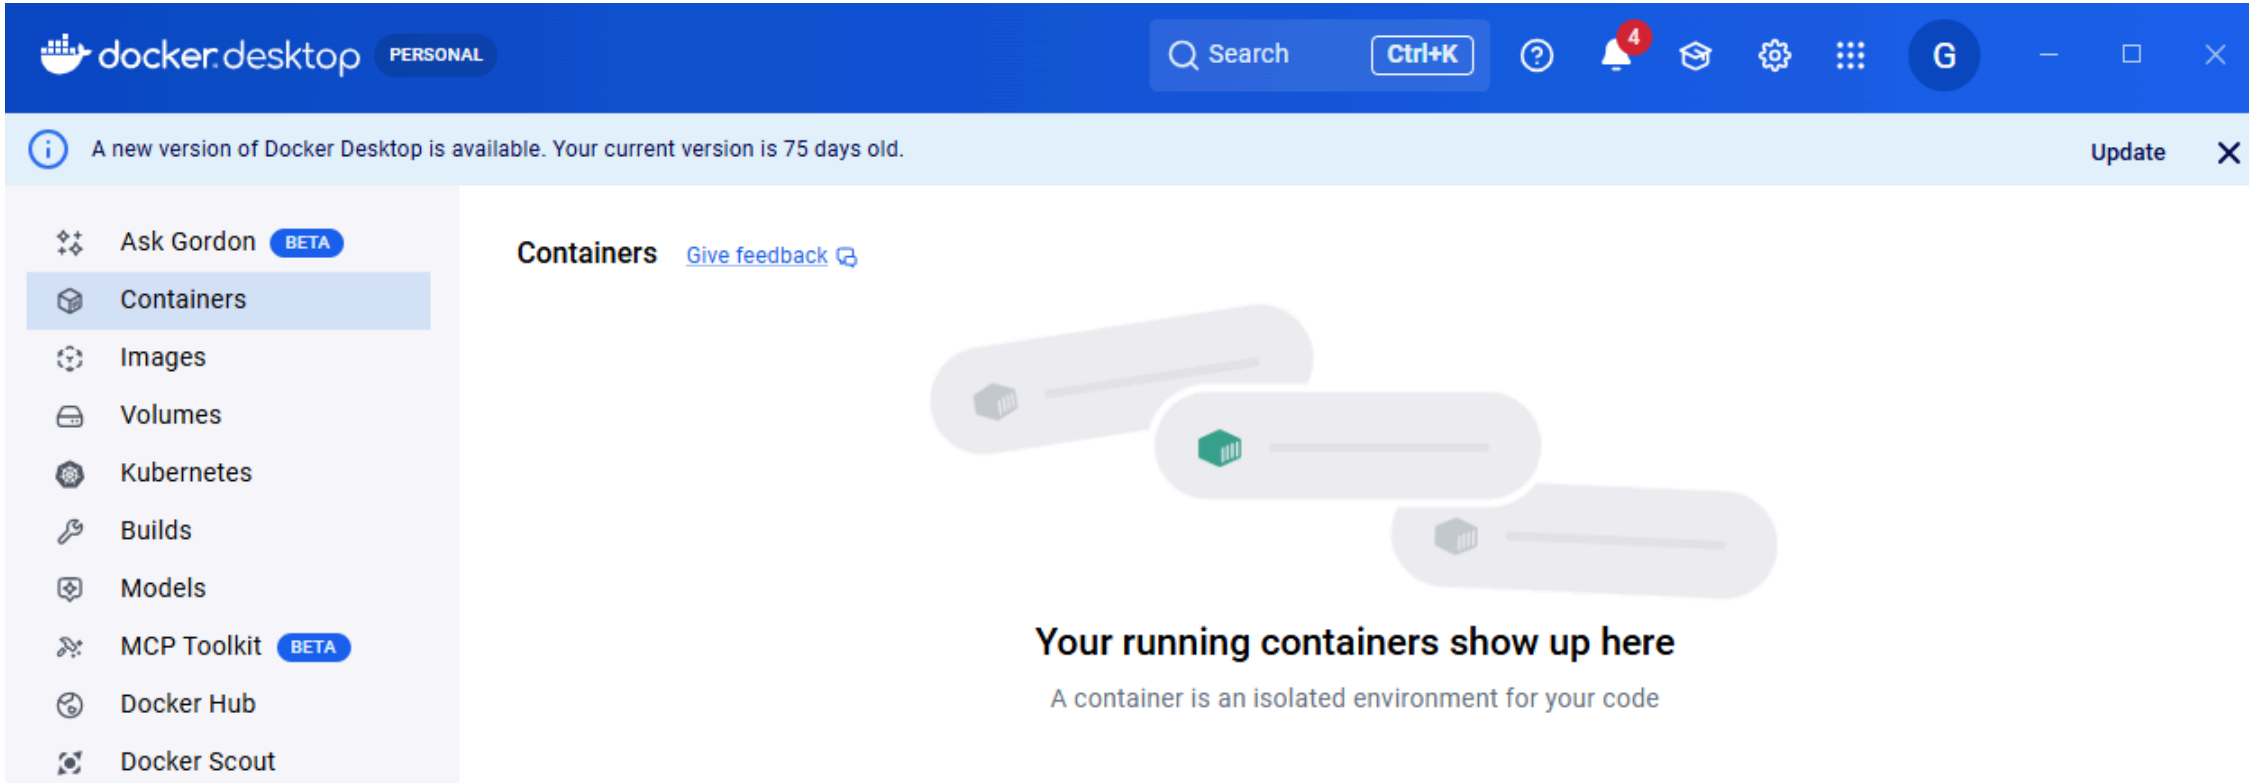

단계2: Docker Image 생성

```
# 도커파일이 있는 폴더에서 실행
docker build --platform linux/amd64 -t [YOUR_USERNAME]/runpod-vllm:latest .
```

```
PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL    PORTS                                                    powershell

PS C:\dev\course\course_LLM\2. LLM Fine Tuning\2. Fine Tuning with Runpod\5. 배포\2-2. runpod with vllm> docker build --platform linux/amd64 -t goodwon593/
runpod-vllm:latest .
[+] Building 2.1s (11/11) FINISHED                                                                                                docker:desktop-linux
 => [internal] load build definition from Dockerfile                                                                                              0.0s
 => => transferring dockerfile: 686B                                                                                                              0.0s
 => [internal] load metadata for docker.io/vllm/vllm-openai:v0.4.2                                                                               1.8s
 => [auth] vllm/vllm-openai:pull token for registry-1.docker.io                                                                                   0.0s
 => [internal] load .dockerignore                                                                                                                 0.0s
 => => transferring context: 2B                                                                                                                   0.0s
 => [1/6] FROM docker.io/vllm/vllm-openai:v0.4.2@sha256:c63bd9e99c6d5914032b396be38a80fd9640693da62b64b907434d38bf38a8e9                          0.0s
 => => resolve docker.io/vllm/vllm-openai:v0.4.2@sha256:c63bd9e99c6d5914032b396be38a80fd9640693da62b64b907434d38bf38a8e9                          0.0s
 => [internal] load build context                                                                                                                 0.0s
 => => transferring context: 1.76kB                                                                                                               0.0s
 => CACHED [2/6] COPY requirements.txt .                                                                                                          0.0s
```

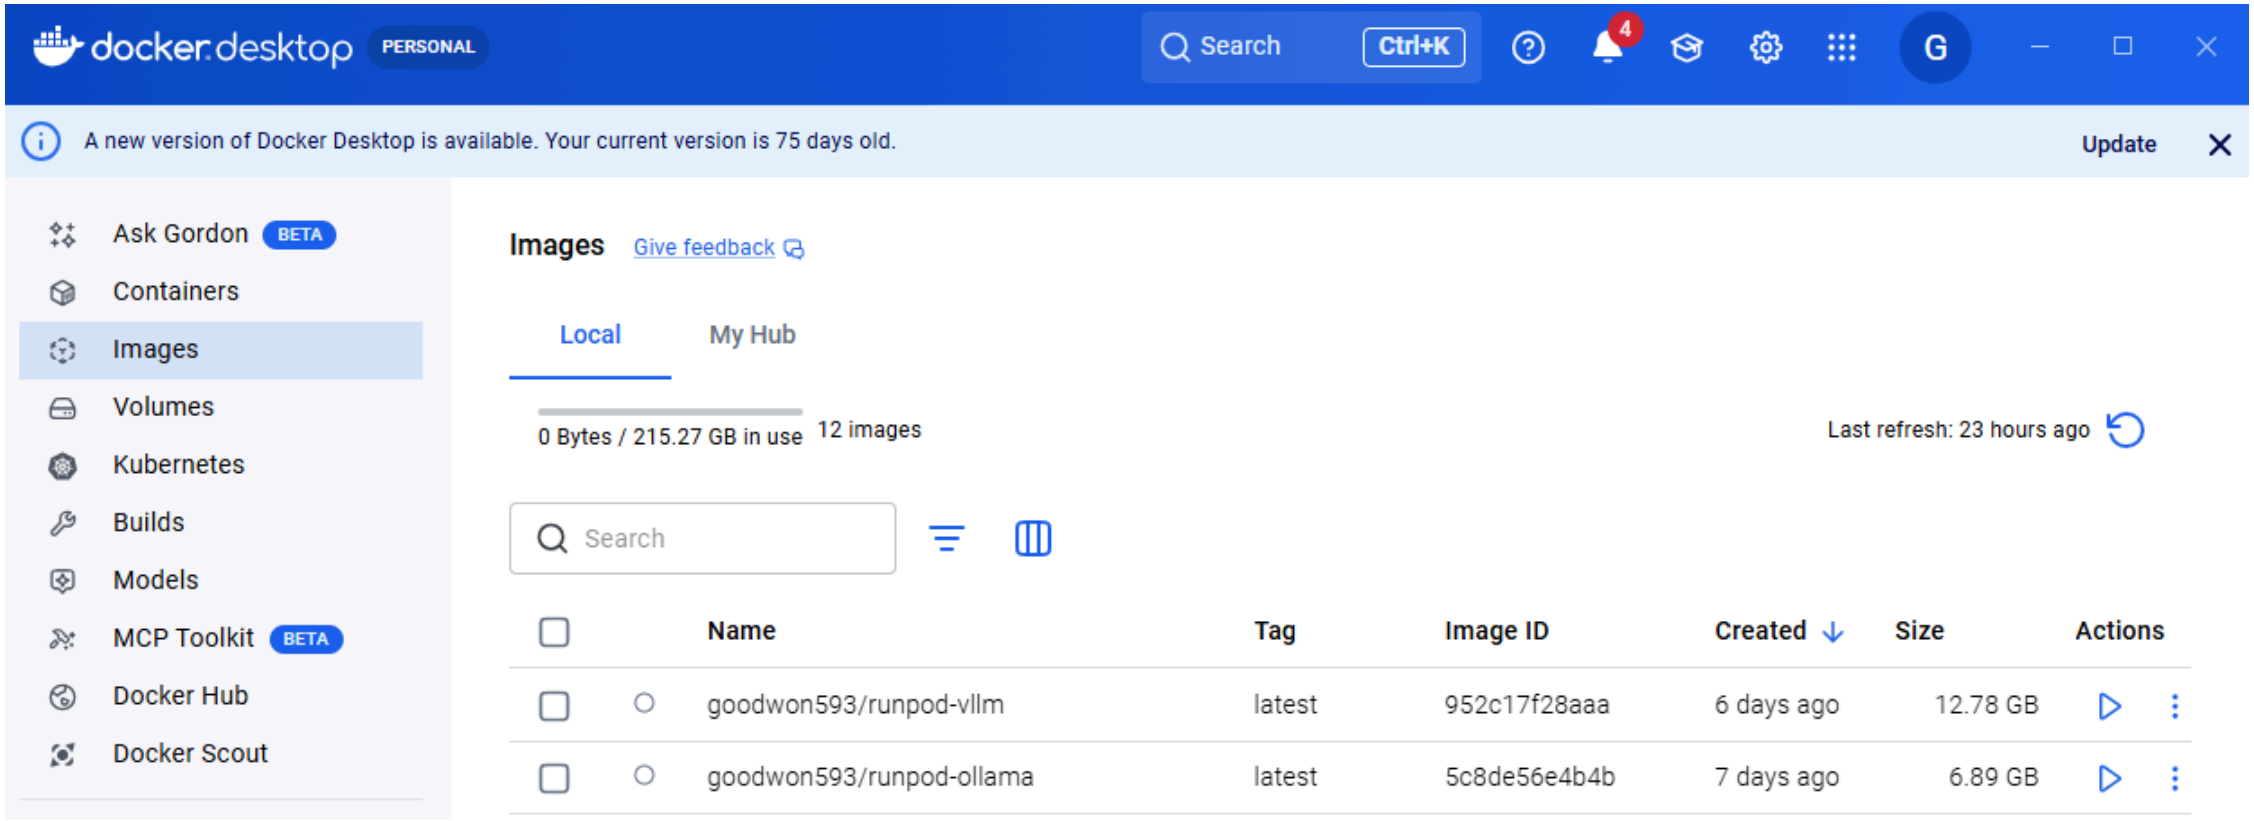
docker desktop
PERSONAL
Search
Ctrl+K
4
G
A new version of Docker Desktop is available. Your current version is 75 days old.
Update
Ask Gordon
BETA
Containers
Images
Volumes
Kubernetes
Builds
Models
MCP Toolkit
BETA
Docker Hub
Docker Scout
Images
Give feedback
Local
My Hub
0 Bytes / 215.27 GB in use
12 images
Last refresh: 23 hours ago
Search
Name
Tag
Image ID
Created
Size
Actions
goodwon593/runpod-vllm
latest
952c17f28aaa
6 days ago
12.78 GB
goodwon593/runpod-ollama
latest
5c8de56e4b4b
7 days ago
6.89 GB

단계3: Docker Hub 배포

```
docker push [YOUR_USERNAME]/runpod-vllm:latest
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS powershell

```
PS C:\dev\course\course_LLM\2. LLM Fine Tuning\2. Fine Tuning with Runpod\5. 배포\2-2. runpod with vllm> docker push goodwon593/runpod-vllm:latest
The push refers to repository [docker.io/goodwon593/runpod-vllm]
b29a99b92af1: Waiting
3b6711cb59bd: Layer already exists
8eca34dd1059: Layer already exists
45f7ea5367fe: Layer already exists
e9c4d5cb6d90: Layer already exists
da5a484f9d74: Layer already exists
0e48f89f4141: Layer already exists
33274e3560b0: Pushed
a6a95425a501: Layer already exists
3d97a47c3c73: Layer already exists
aece8493d397: Layer already exists
```

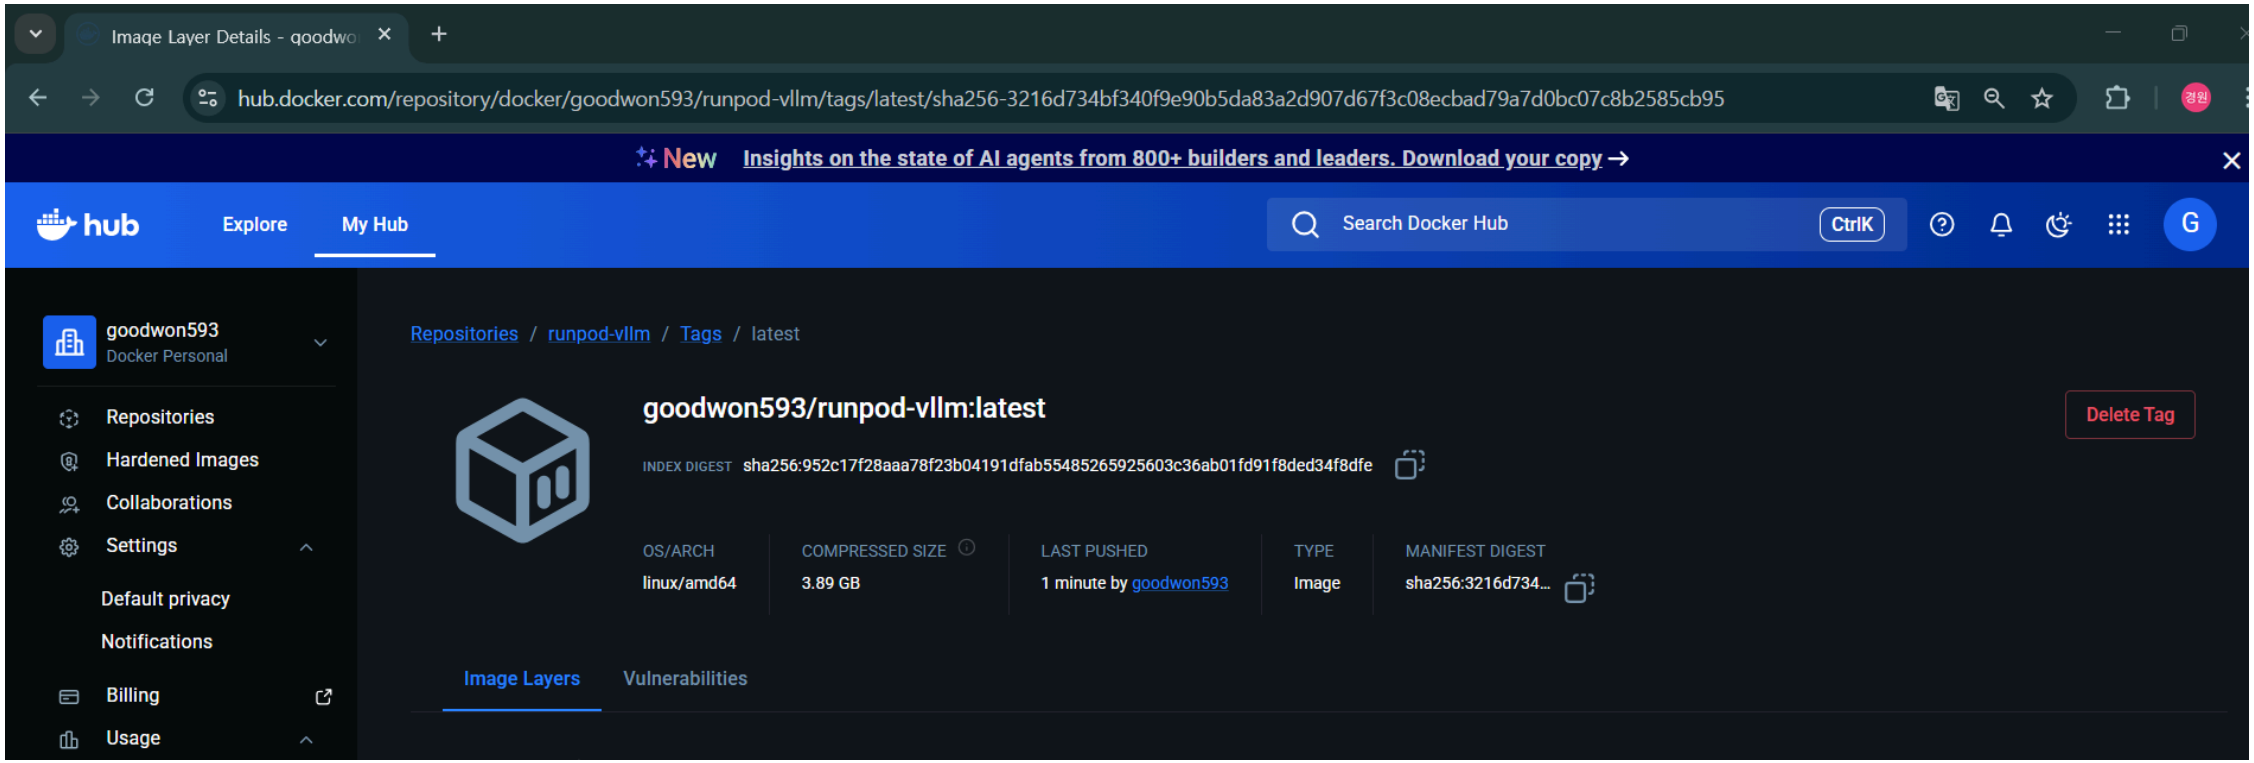
Image Layer Details - goodwo
hub.docker.com/repository/docker/goodwon593/runpod-vllm/tags/latest/sha256-3216d734bf340f9e90b5da83a2d907d67f3c08ecbad79a7d0bc07c8b2585cb95
New
Insights on the state of AI agents from 800+ builders and leaders. Download your copy →
hub
Explore
My Hub
Search Docker Hub
CtrlK
G
goodwon593
Docker Personal
Repositories
Hardened Images
Collaborations
Settings
Default privacy
Notifications
Billing
Usage
Repositories / runpod-vllm / Tags / latest
goodwon593/runpod-vllm:latest
Delete Tag
INDEX DIGEST sha256:952c17f28aaa78f23b04191dfab55485265925603c36ab01fd91f8ded34f8dfe
OS/ARCH
linux/amd64
COMPRESSED SIZE
3.89 GB
LAST PUSHED
1 minute by goodwon593
TYPE
Image
MANIFEST DIGEST
sha256:3216d734...
Image Layers
Vulnerabilities

Runpod

단계 1: Template 생성

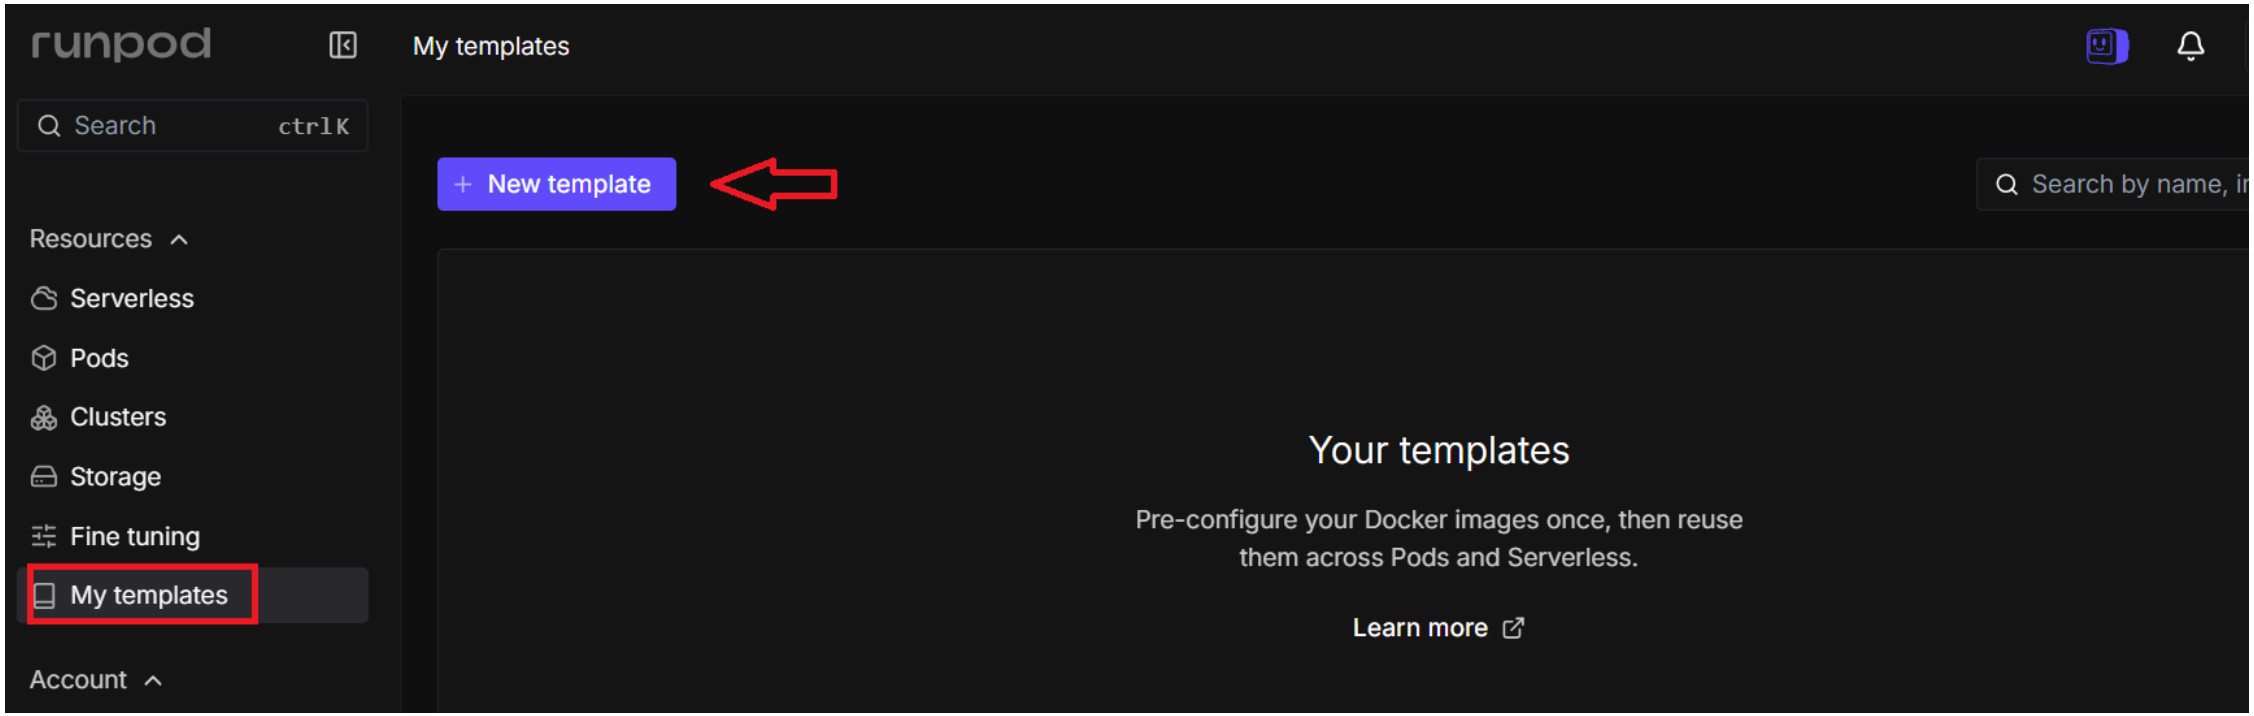

Name, Public template 등

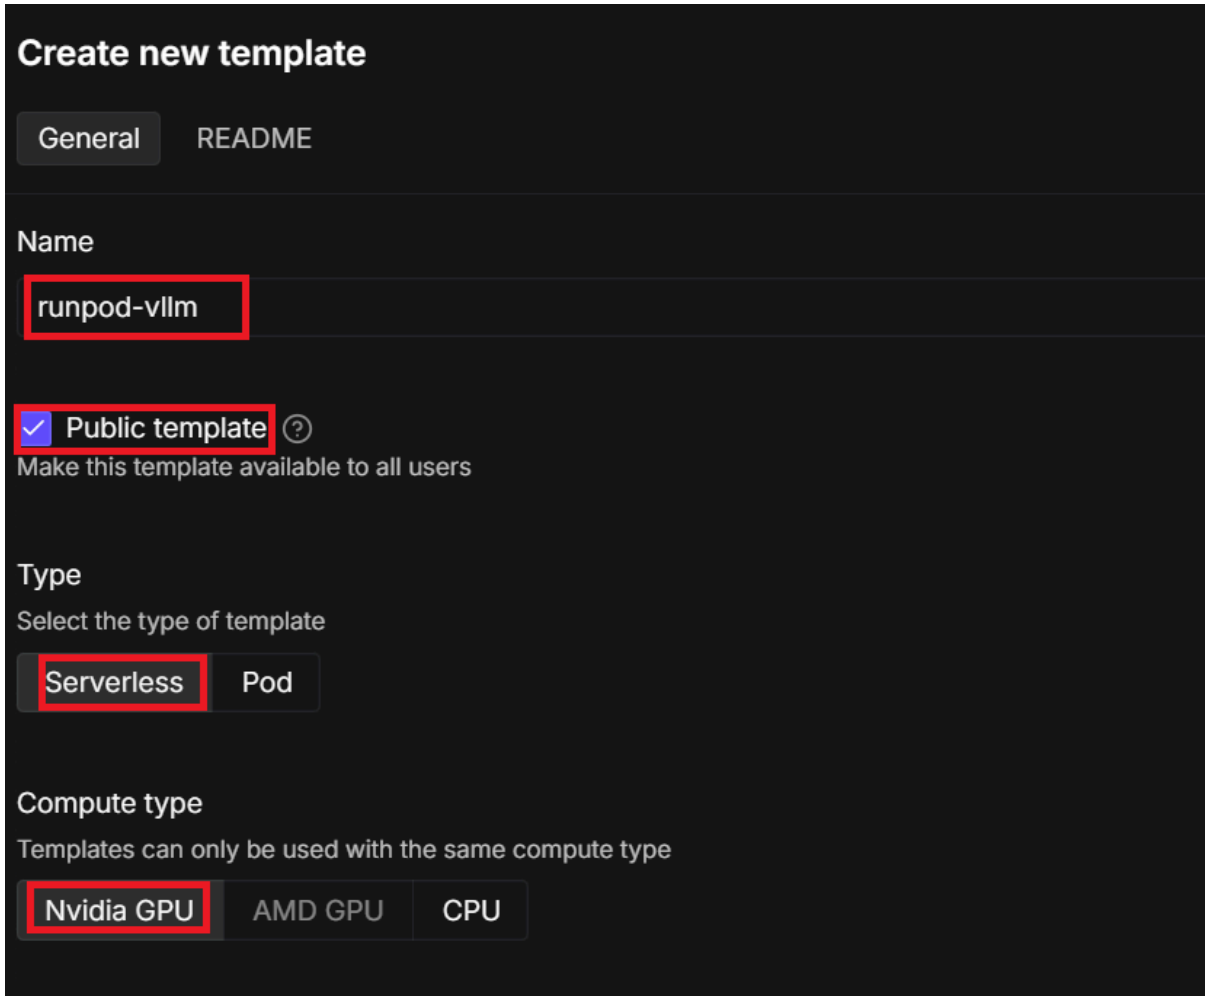

Container image 등

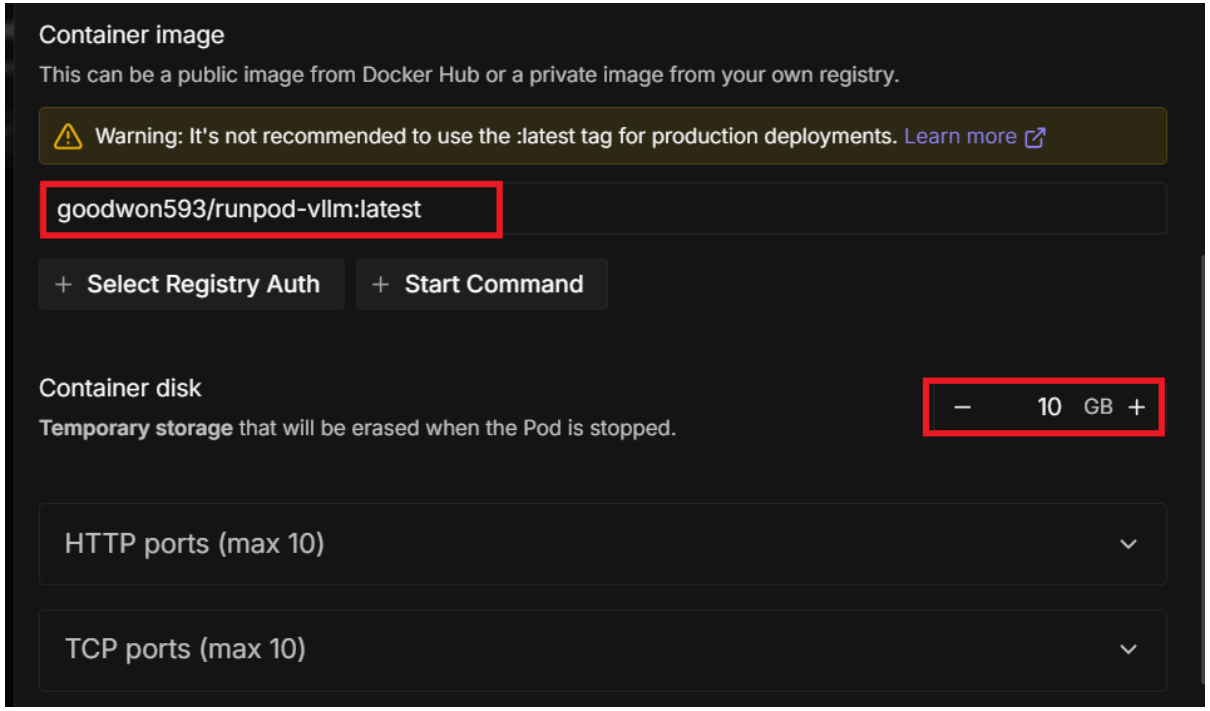

(옵션) 환경변수 등록

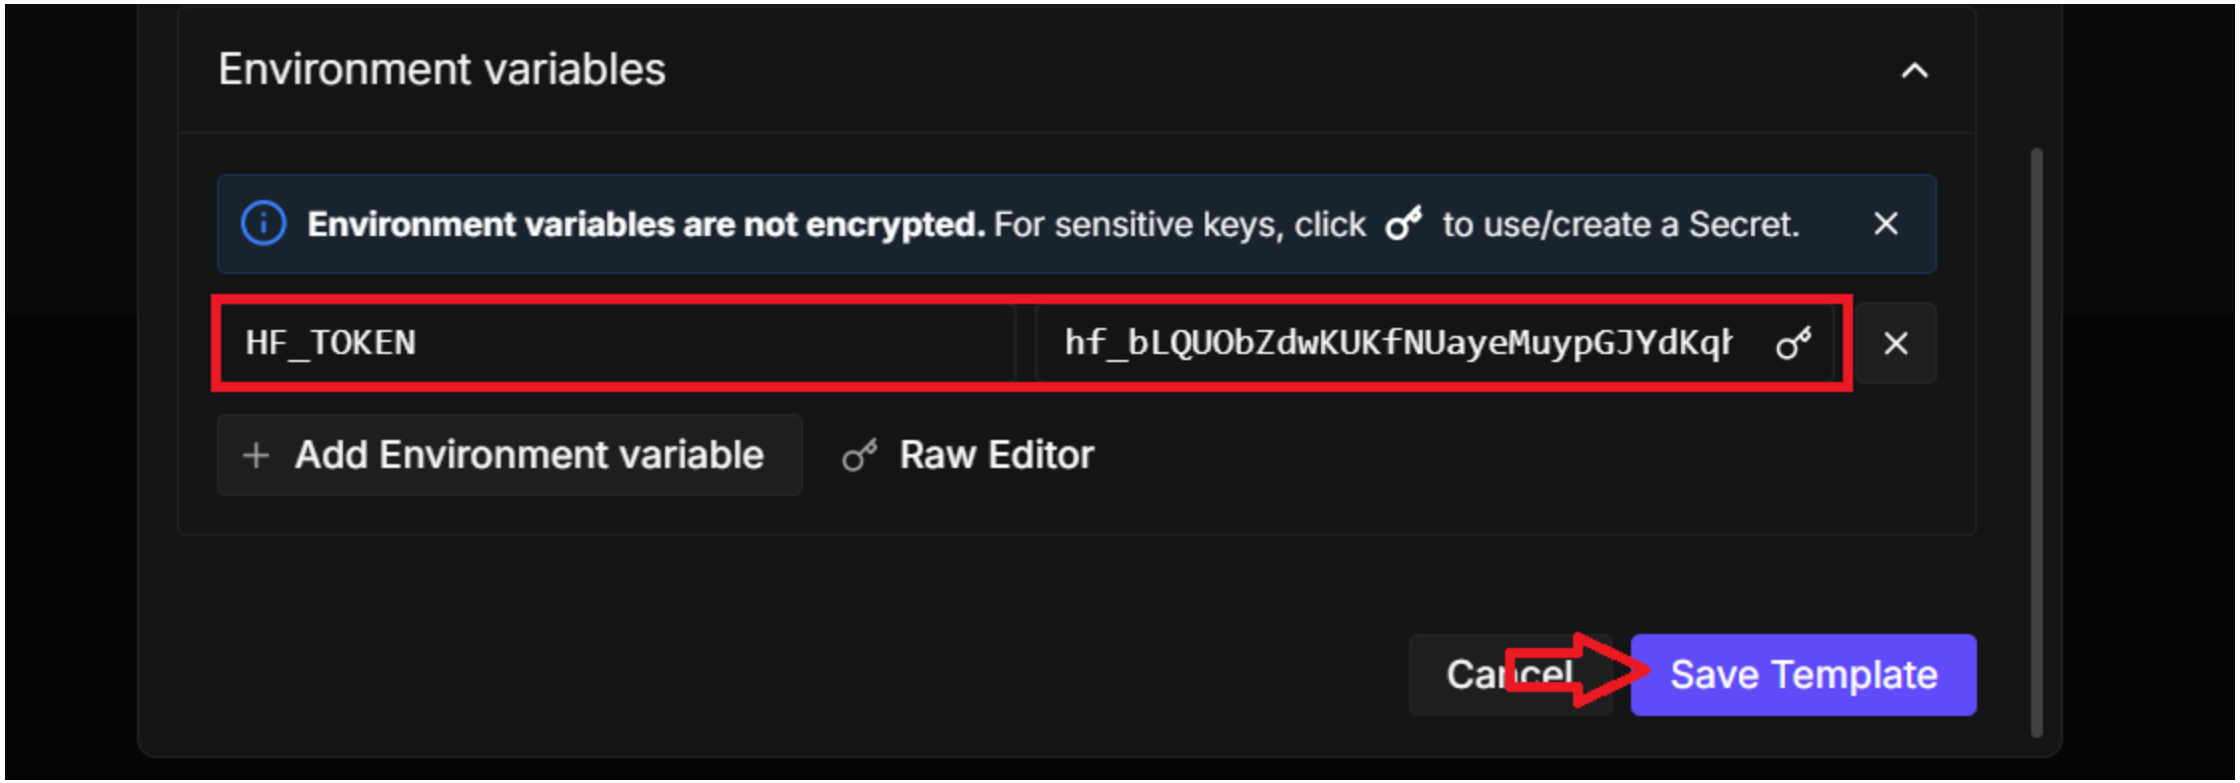
Environment variables
Environment variables are not encrypted. For sensitive keys, click to use/create a Secret.
HF_TOKEN
hf_bLQUObZdwKUKfNUayeMuypGJYdKqh
Add Environment variable
Raw Editor
Cancel
Save Template

결과 확인

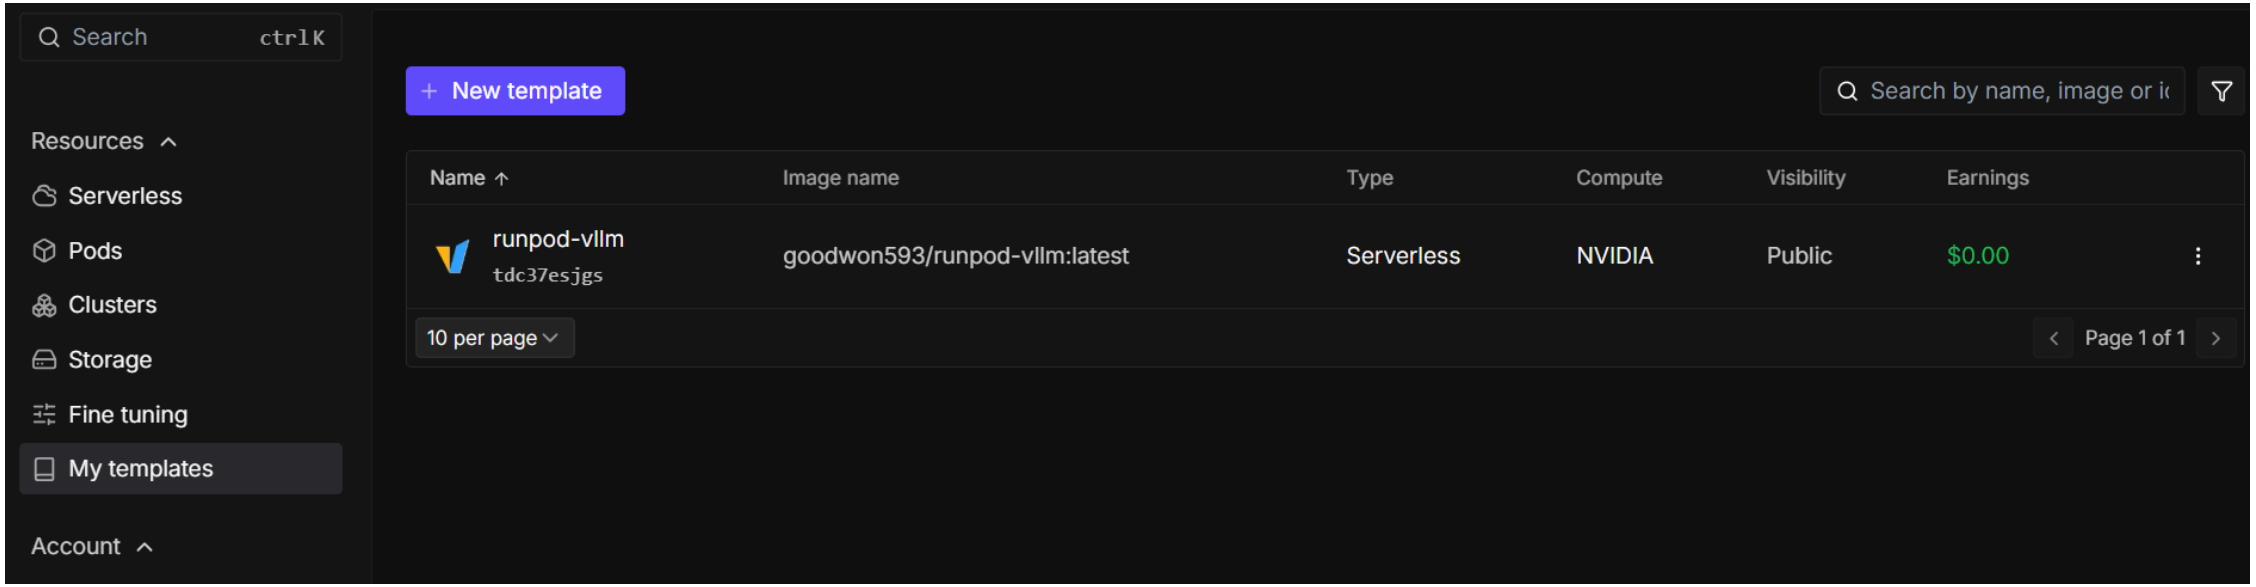
Search
ctrlK
Resources
Serverless
Pods
Clusters
Storage
Fine tuning
My templates
Account
New template
Search by name, image or i
Name
Image name
Type
Compute
Visibility
Earnings
runpod-vllm
tdc37esjgs
goodwon593/runpod-vllm:latest
Serverless
NVIDIA
Public
$0.00
10 per page
Page 1 of 1

단계 2: Serverless 생성

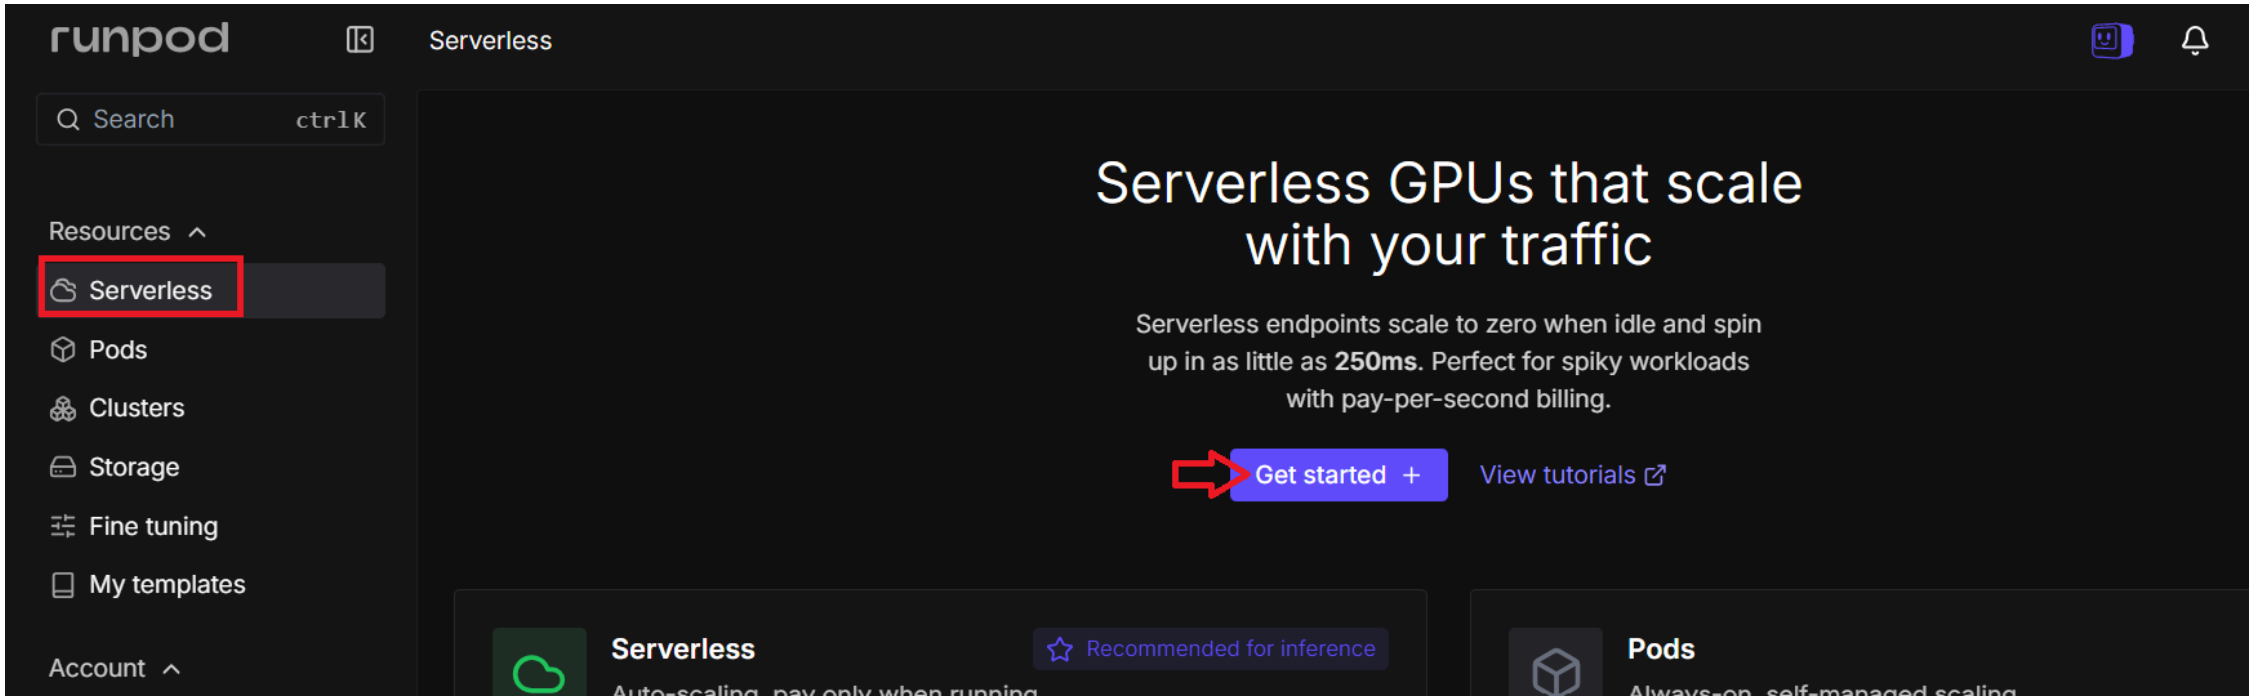

Choose a template

Resources
- Serverless
- Pods
- Clusters
- Storage
- Fine tuning
- My templates

Account
- Savings plans
- Audit logs
- Remote access
- Refer & earn

Feedback

Help & resources

Create a new deployment

Not sure which to choose? Ask the Runpod Assistant

Start from a Hub listing
Choose from pre-built models and applications. No code or configuration required.

Run code locally Flash Beta
Write Python locally in your IDE, execute remotely on GPUs. Just add @Endpoint.

Custom deployment
Configure a deployment from scratch with your own Docker image, GitHub repo, or custom template.

Deploy a custom configuration

Deploy from GitHub
Deploy directly from your Github repository

Deploy from Docker registry or a template
Deploy from Dockerhub, other image repositories, or a custom template

Choose a template
Select a template to use for this endpoint

runpod-vllm

You have selected a template. Clear the template to edit values directly.

Configure endpoint

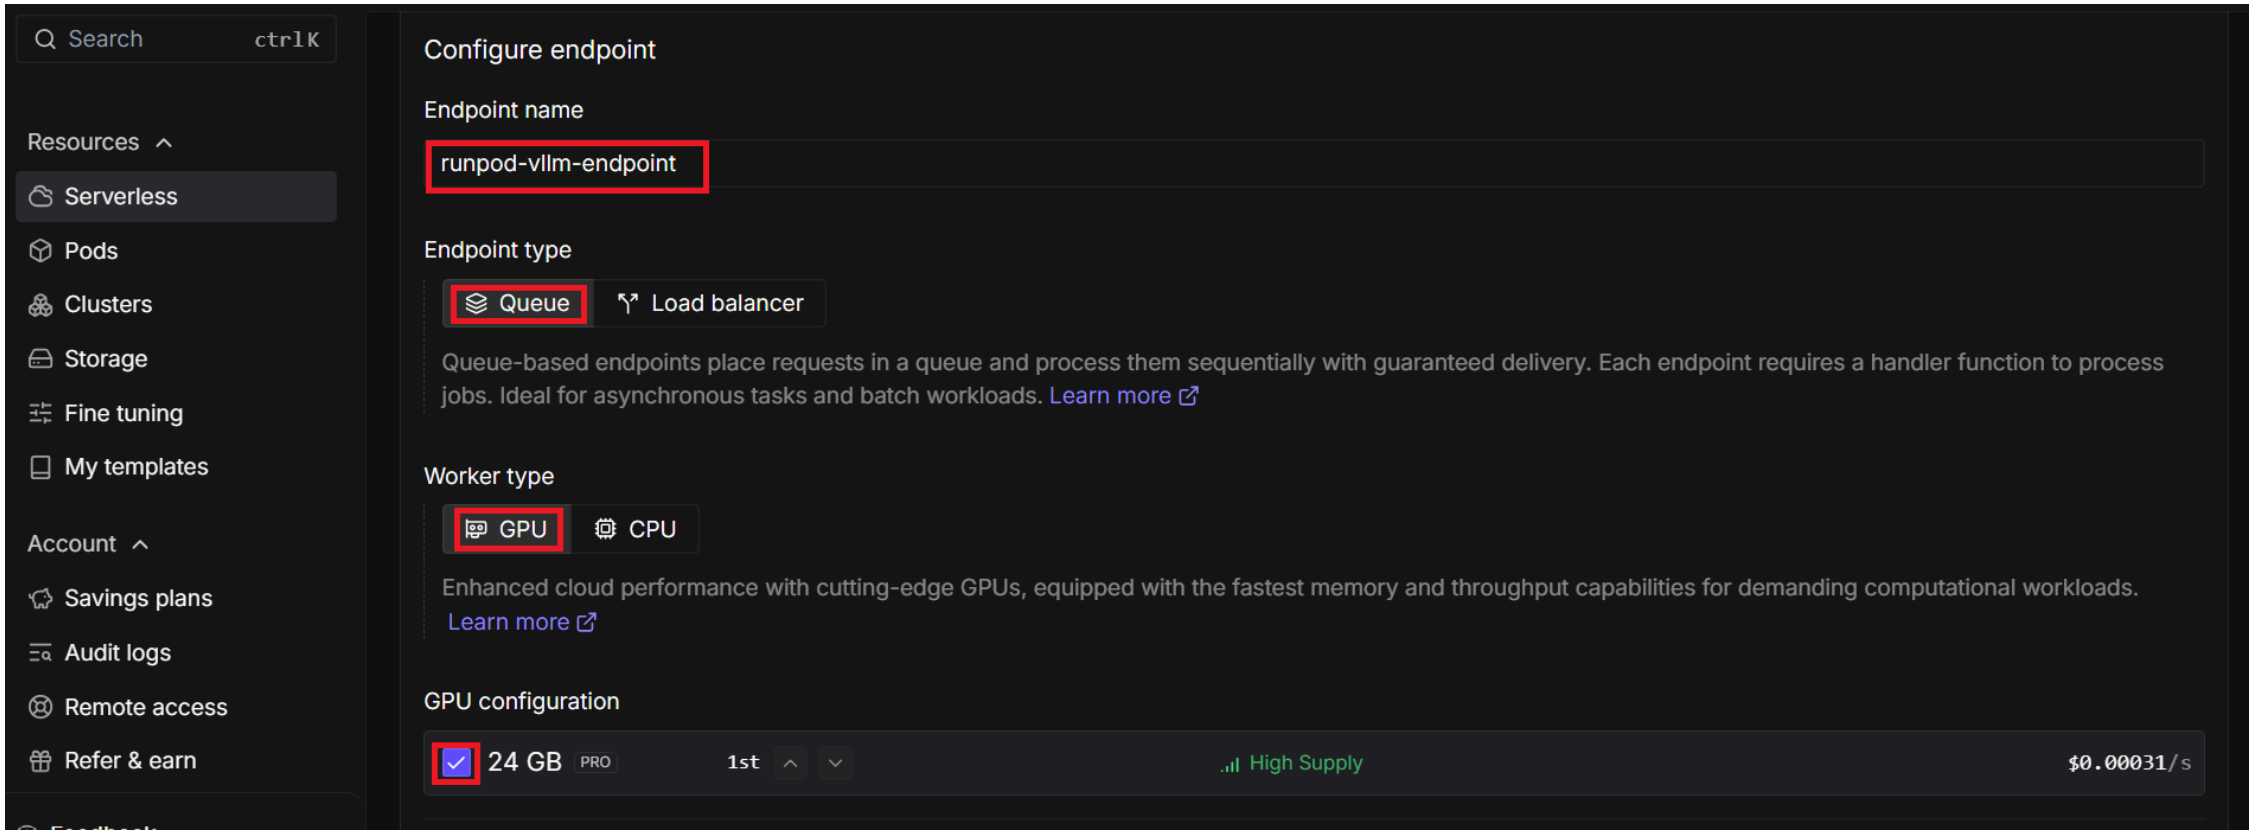

Create endpoint

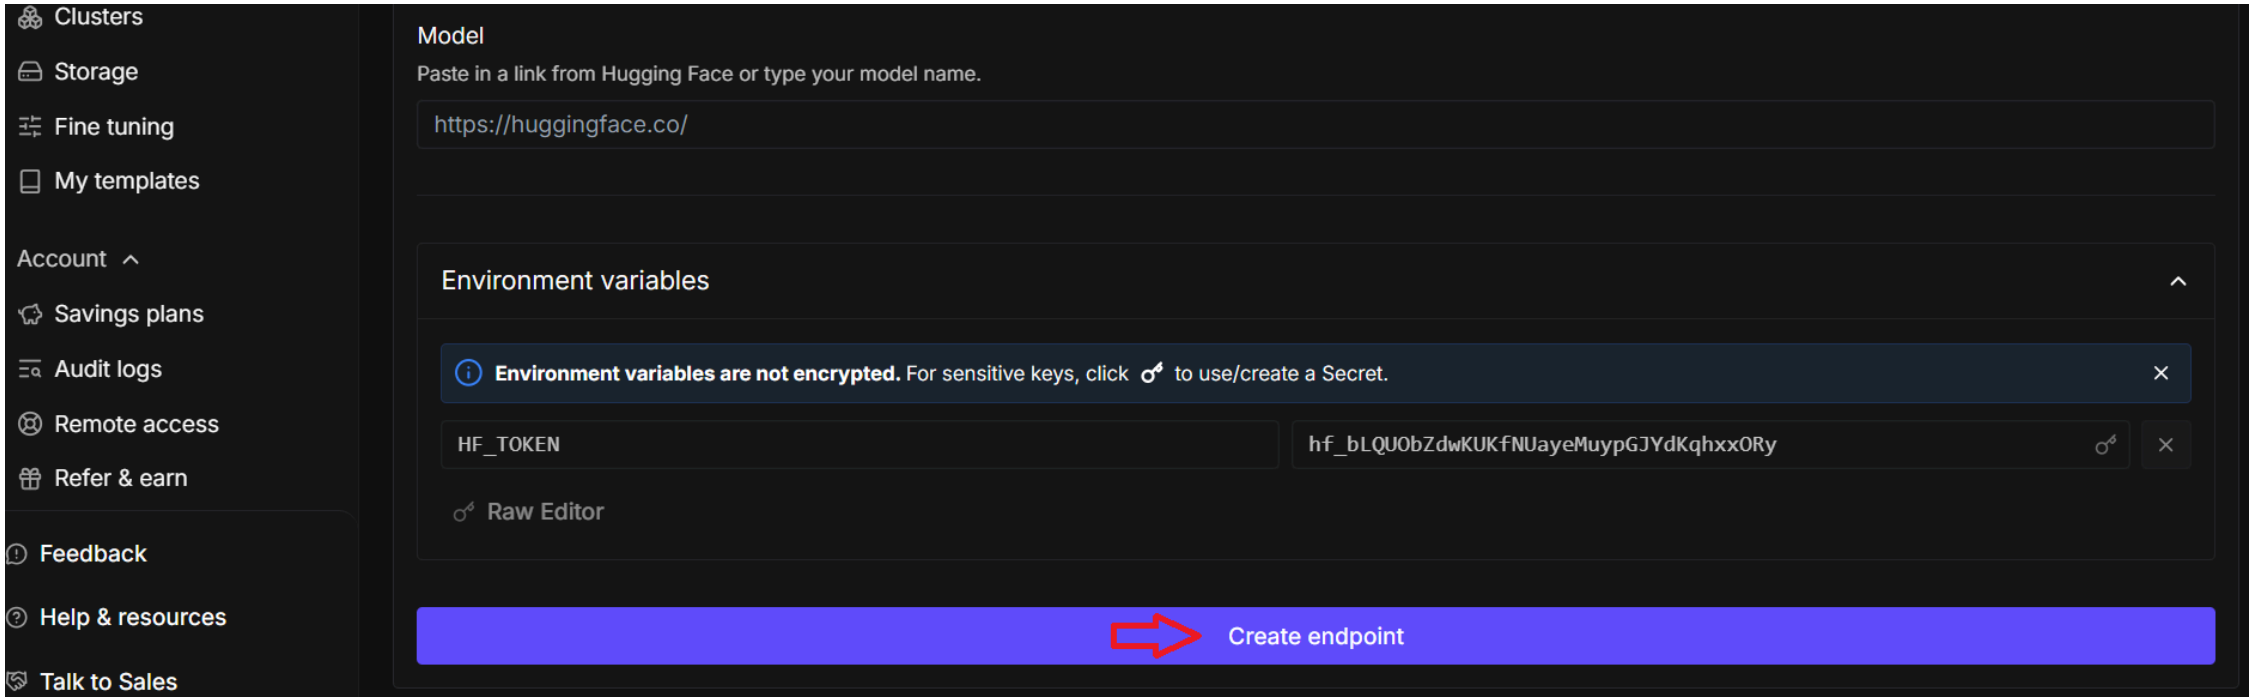

단계 3: Active Workers 적용

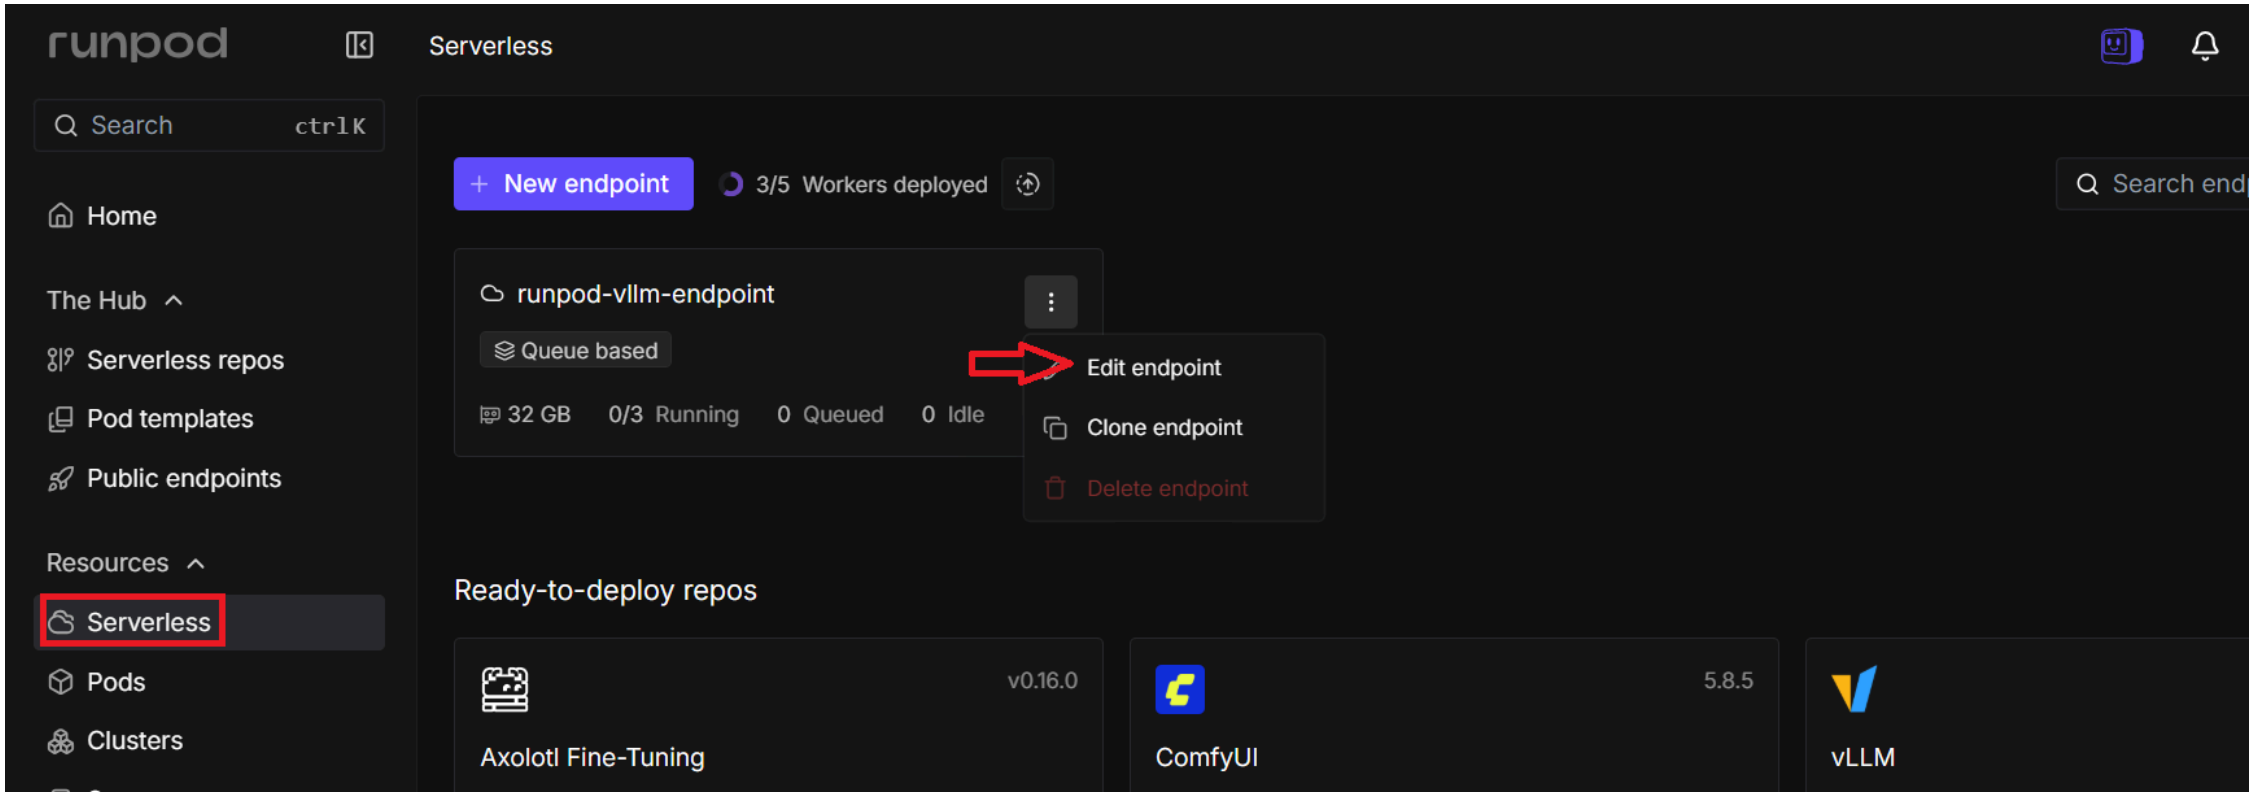

Active Workers

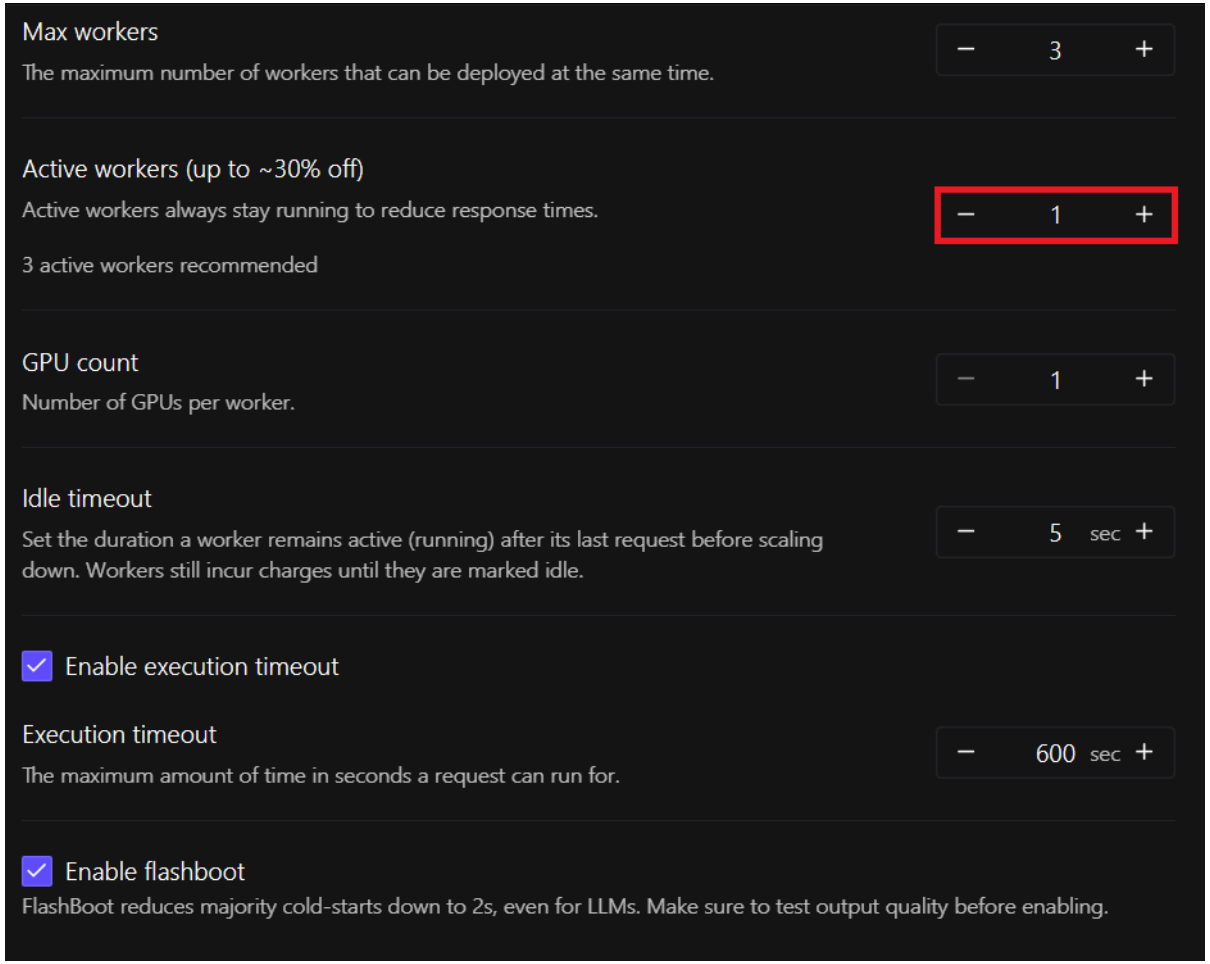

Save Endpoint

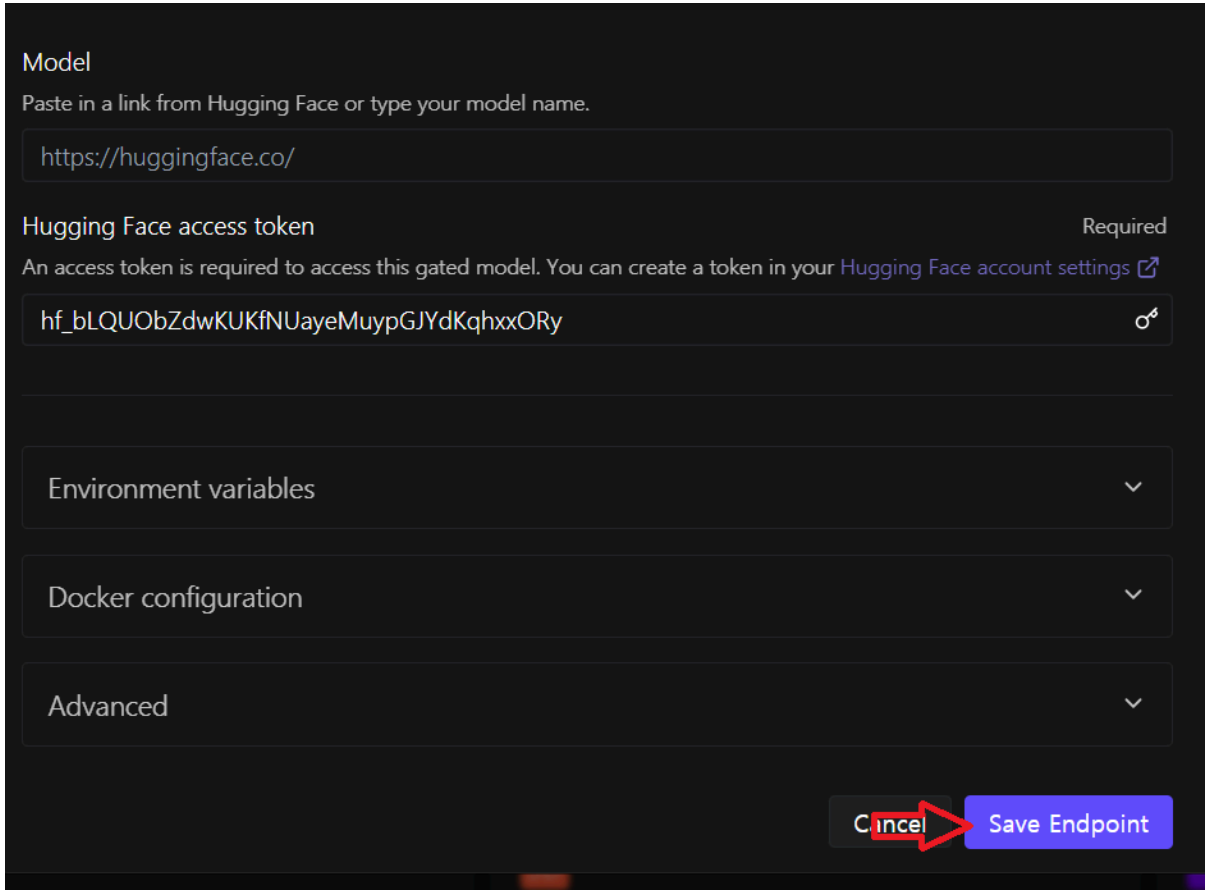

적용 확인

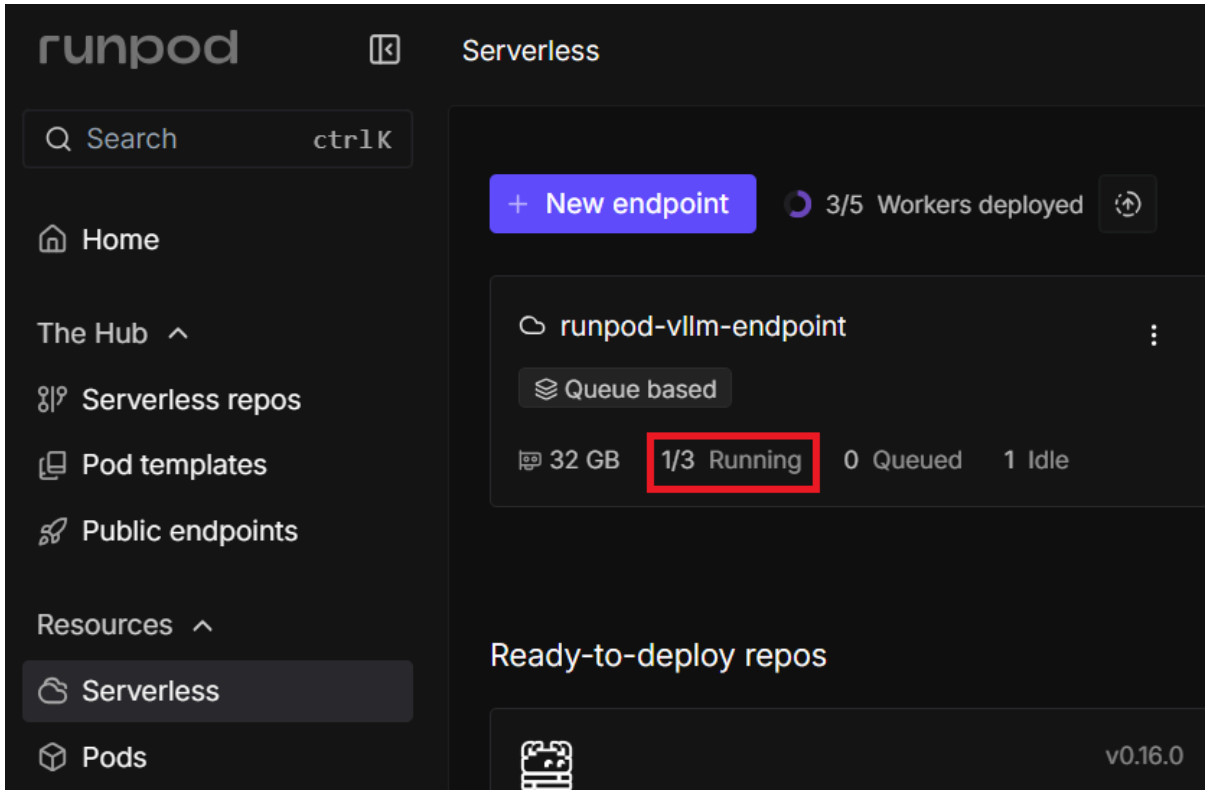
runpod
Serverless
Search
ctrlK
Home
The Hub
Serverless repos
Pod templates
Public endpoints
Resources
Serverless
Pods
+ New endpoint
3/5 Workers deployed
runpod-vllm-endpoint
Queue based
32 GB
1/3 Running
0 Queued
1 Idle
Ready-to-deploy repos
v0.16.0

단계 4: 테스트

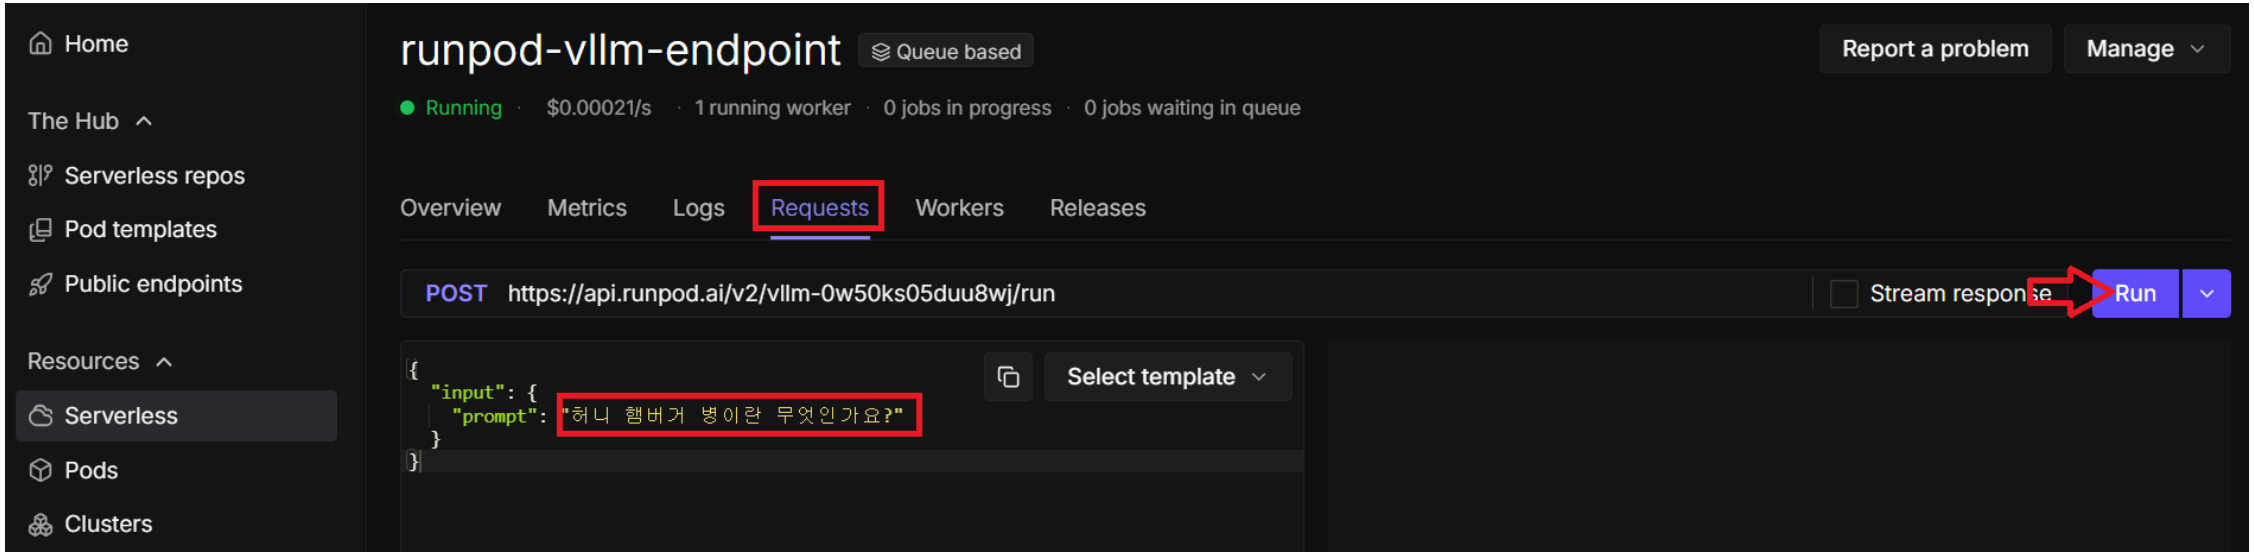
Home
The Hub
Serverless repos
Pod templates
Public endpoints
Resources
Serverless
Pods
Clusters
runpod-vllm-endpoint
Queue based
Report a problem
Manage
Running · $0.00021/s · 1 running worker · 0 jobs in progress · 0 jobs waiting in queue
Overview
Metrics
Logs
Requests
Workers
Releases
POST https://api.runpod.ai/v2/vllm-0w50ks05duu8wj/run
Stream response
Run
Select template
{
"input": {
"prompt": "허니 햄버거 병이란 무엇인가요?"
}
}

실행 중인 Worker 확인

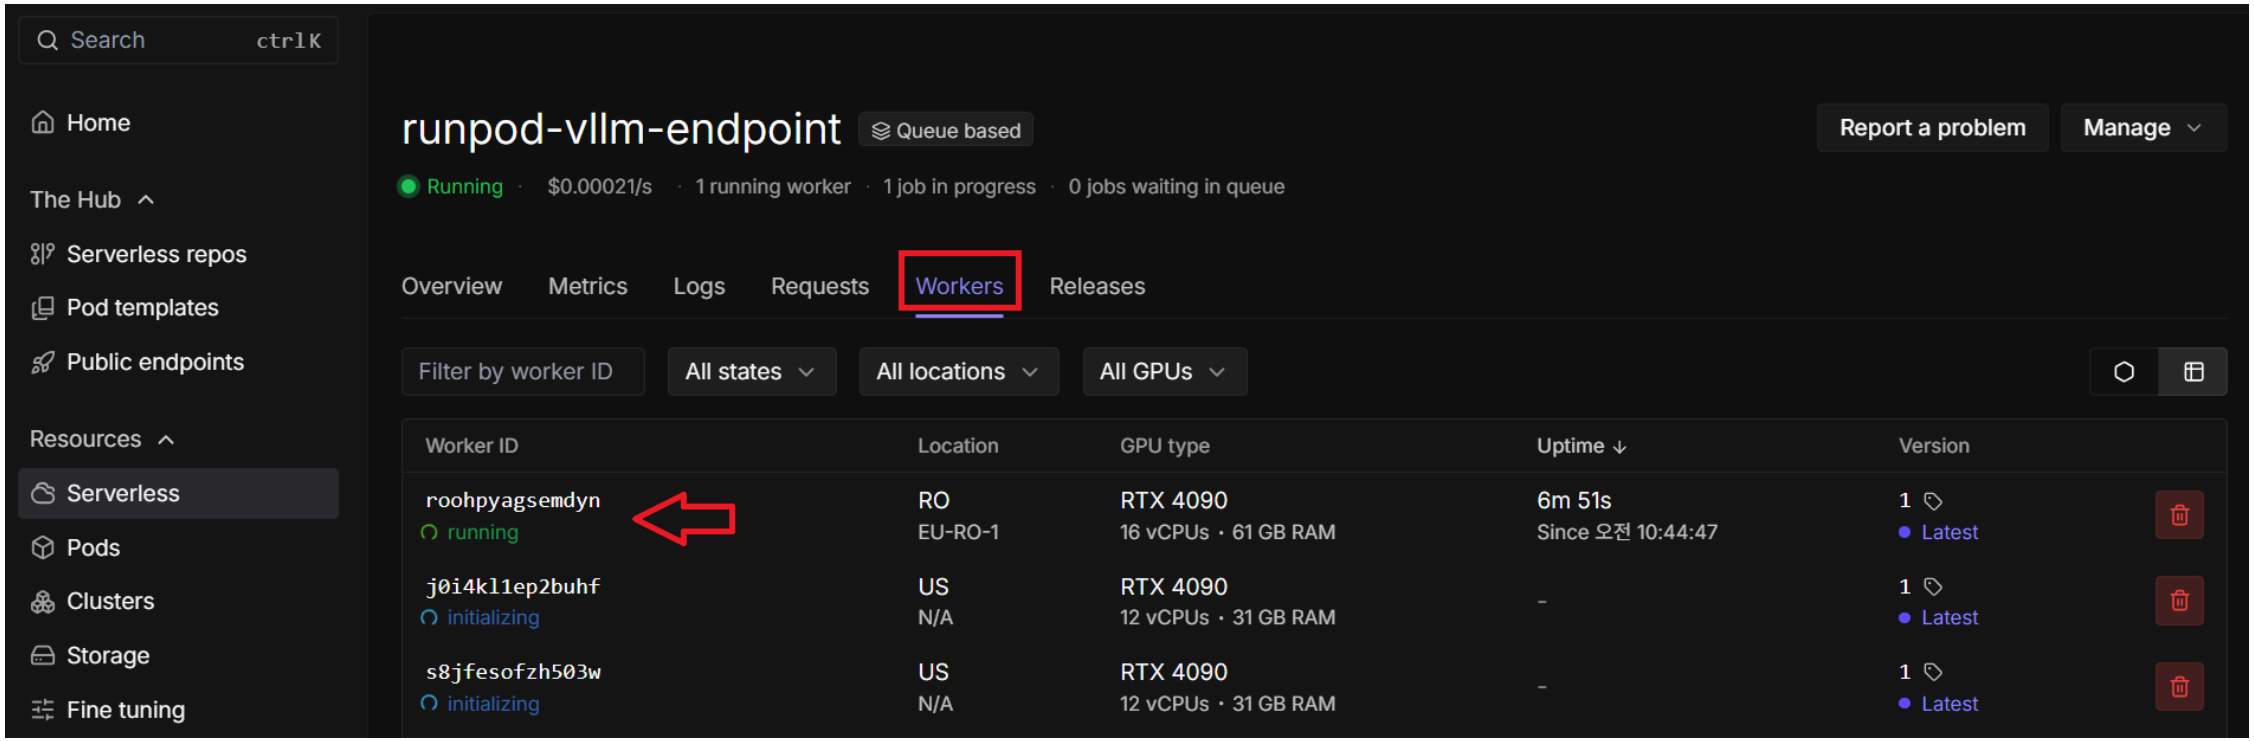

log 확인

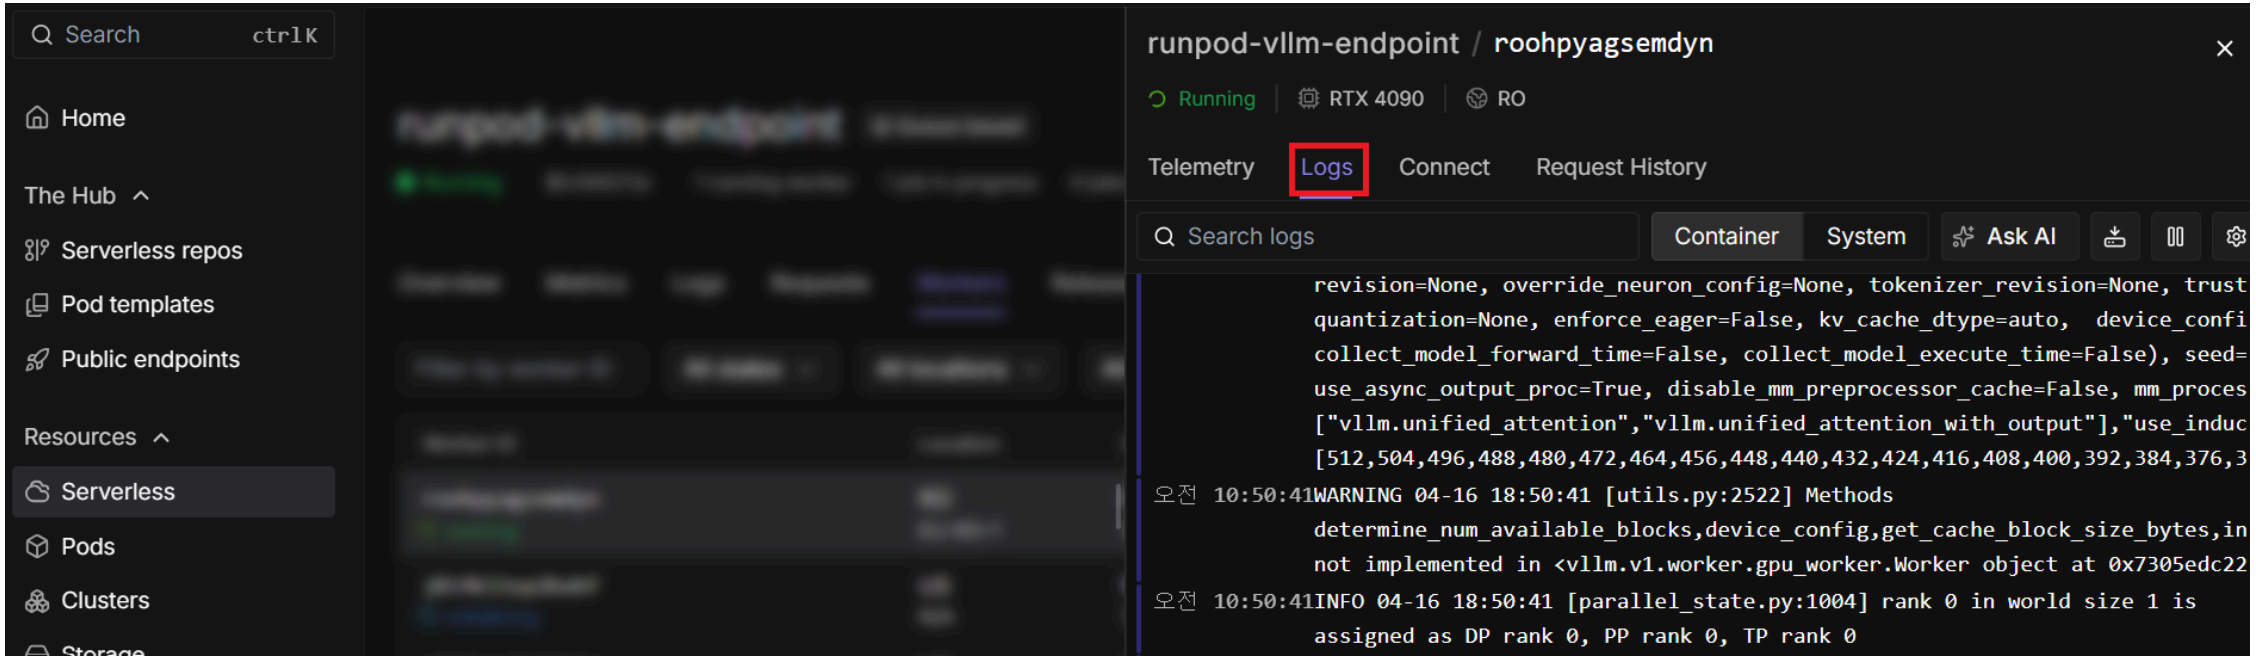
Search ctrlK
Home
The Hub
Serverless repos
Pod templates
Public endpoints
Resources
Serverless
Pods
Clusters
runpod-vllm-endpoint / roohpyagsemdyn
Running
RTX 4090
RO
Telemetry
Logs
Connect
Request History
Search logs
Container
System
Ask AI
revision=None, override_neuron_config=None, tokenizer_revision=None, trust
quantization=None, enforce_eager=False, kv_cache_dtype=auto, device_confi
collect_model_forward_time=False, collect_model_execute_time=False), seed=
use_async_output_proc=True, disable_mm_preprocessor_cache=False, mm_proces
["vllm.unified_attention","vllm.unified_attention_with_output"],"use_induc
[512,504,496,488,480,472,464,456,448,440,432,424,416,408,400,392,384,376,3
오전 10:50:41WARNING 04-16 18:50:41 [utils.py:2522] Methods
determine_num_available_blocks,device_config,get_cache_block_size_bytes,in
not implemented in <vllm.v1.worker.gpu_worker.Worker object at 0x7305edc22
오전 10:50:41INFO 04-16 18:50:41 [parallel_state.py:1004] rank 0 in world size 1 is
assigned as DP rank 0, PP rank 0, TP rank 0

결과 확인

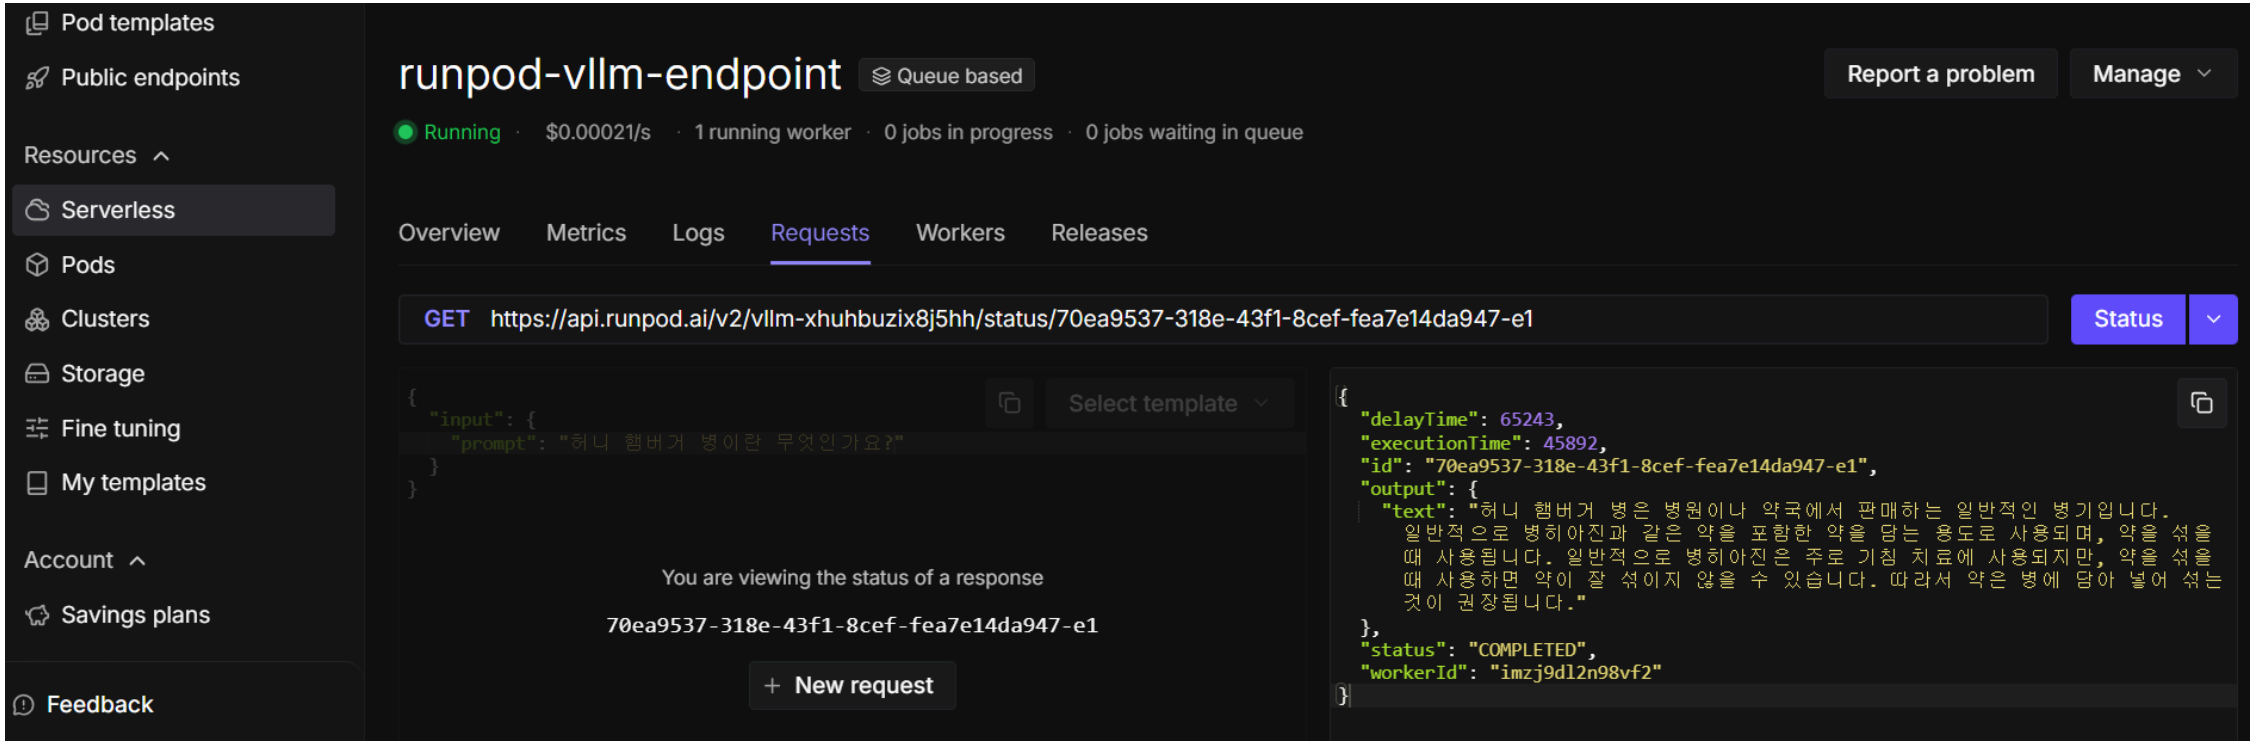
Pod templates
Public endpoints
Resources
Serverless
Pods
Clusters
Storage
Fine tuning
My templates
Account
Savings plans
Feedback
runpod-vllm-endpoint
Queue based
Report a problem
Manage
Running · $0.00021/s · 1 running worker · 0 jobs in progress · 0 jobs waiting in queue
Overview
Metrics
Logs
Requests
Workers
Releases
GET https://api.runpod.ai/v2/vllm-xhuhbuzix8j5hh/status/70ea9537-318e-43f1-8cef-fea7e14da947-e1
Status
Select template
You are viewing the status of a response
70ea9537-318e-43f1-8cef-fea7e14da947-e1
New request
{
"delayTime": 65243,
"executionTime": 45892,
"id": "70ea9537-318e-43f1-8cef-fea7e14da947-e1",
"output": {
"text": "허니 햄버거 병은 병원이나 약국에서 판매하는 일반적인 병기입니다. 일반적으로 병히아진과 같은 약을 포함한 약을 담는 용도로 사용되며, 약을 섞을 때 사용됩니다. 일반적으로 병히아진은 주로 기침 치료에 사용되지만, 약을 섞을 때 사용하면 약이 잘 섞이지 않을 수 있습니다. 따라서 약은 병에 담아 넣어 섞는 것이 권장됩니다."
},
"status": "COMPLETED",
"workerId": "imzj9dl2n98vf2"
}